

Deep Learning

5. Overview of classical computer vision - 1.
Basic image processing, filtering, Fourier transform, image
deblurring, edge detection

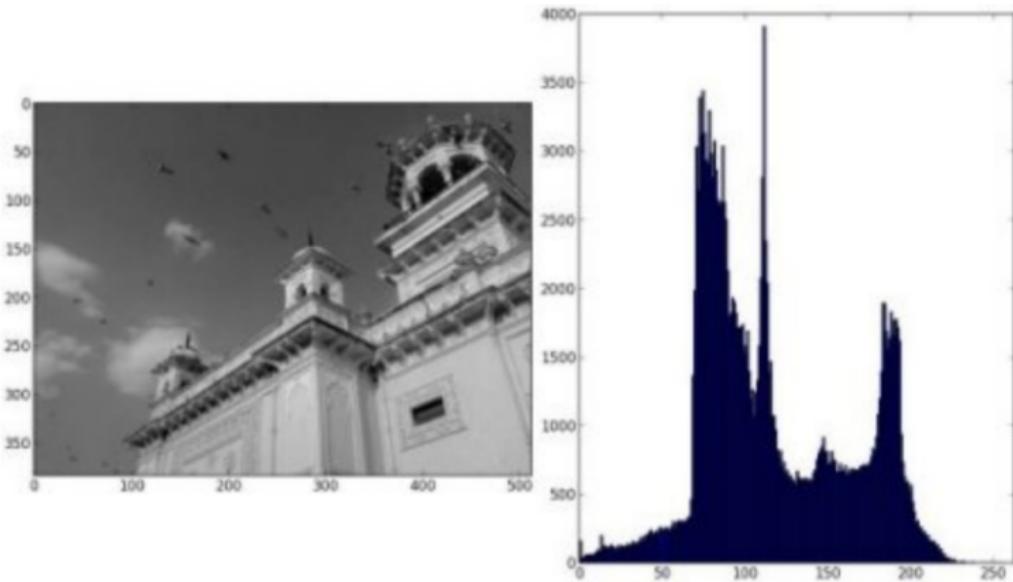
Viacheslav Dudar

Taras Shevchenko National University of Kyiv

2018

Histogram of the image

Histogram of the image: shows distribution of pixel intensities.



Point operations

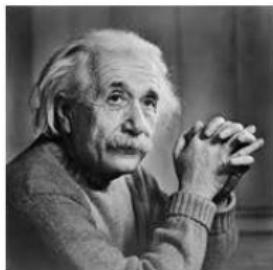
Apply pixelwise transformation to the input image:

$$J(x, y) = F(I(x, y))$$

Example 1: digital negative

$$F(x) = 1 - x$$

Input Image

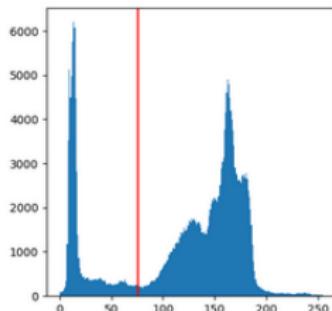
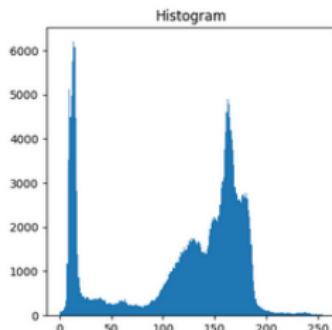


Output Image



Point operations: thresholding

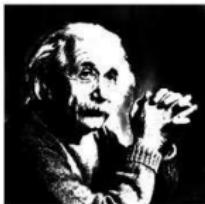
Example 2: image thresholding: $F(x) = (x > \theta)$



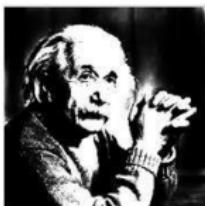
Point operations: power-law

Example 3: Gamma correction: $F(x) = x^\gamma$

Gamma = 10



Gamma = 8



Gamma = 6

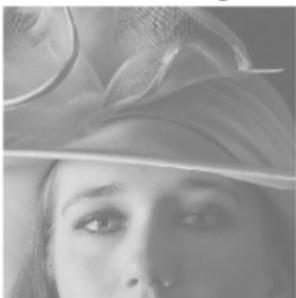


Point operations: contrast stretching

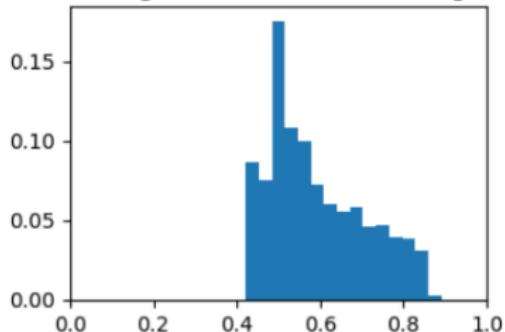
$$g = \frac{f - f_{\min}}{f_{\max} - f_{\min}}$$

Example 4: contrast stretching:

Low contrast orginal



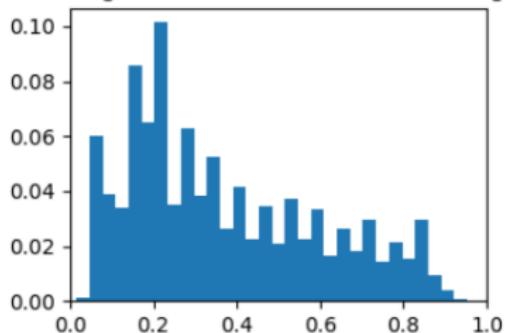
Histogram of low contrast image



Contrast Stretched



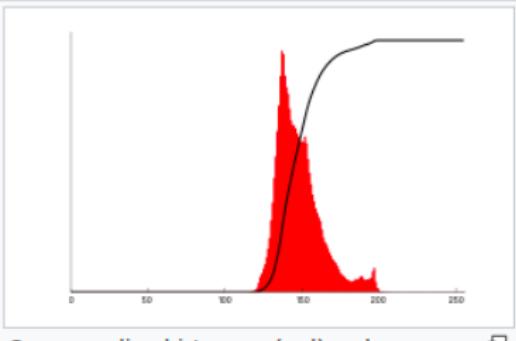
Histogram of contrast stretched image



Point operations: Histogram equalization



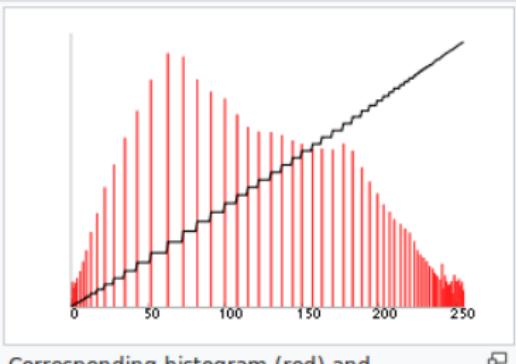
Before Histogram Equalization



Corresponding histogram (red) and
cumulative histogram (black)



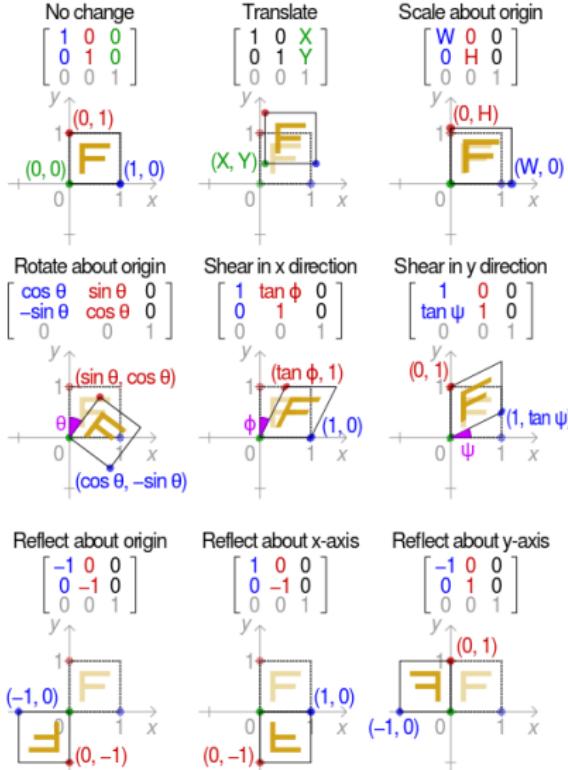
After Histogram Equalization



Corresponding histogram (red) and
cumulative histogram (black)

Geometric transformations

General formula for geometric transformation: $J(x, y) = I(T(x, y))$



Affine transformation:

$$T(x, y) = \begin{pmatrix} ax + by + c \\ dx + ey + f \end{pmatrix}$$

Combination of

- scaling
- rotation
- shear
- translation

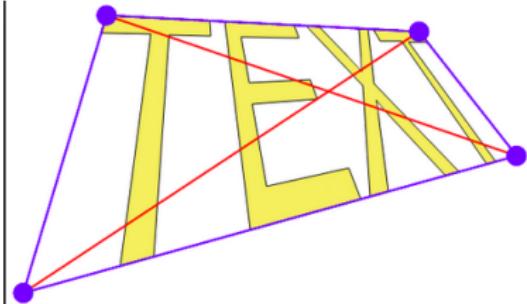
Inverse is also affine

Parallel lines remain parallel

Can convert any to any triangle

Needs 3 points to be found

Geometric transformations



Perspective transformation:

$$T(x, y) = \begin{pmatrix} ax+by+c \\ gx+hy+1 \\ \frac{dx+ey+f}{gx+hy+1} \end{pmatrix}$$

Inverse is also perspective

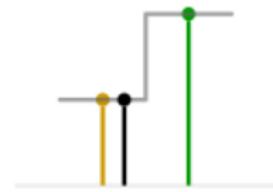
Can convert any to any quadrangle

Needs 4 points to be found

Example: QR codes scan



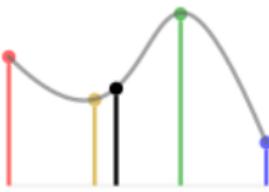
Image interpolations



1D nearest-neighbour



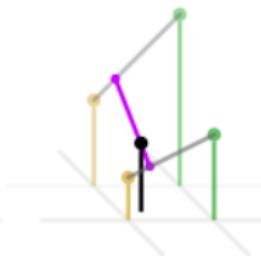
Linear



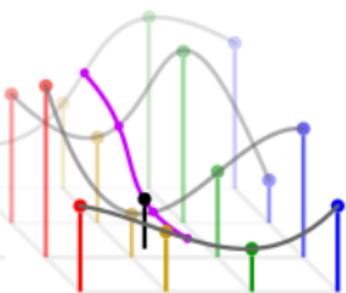
Cubic



2D nearest-neighbour



Bilinear



Bicubic

2D convolution

Image: $I[i, j]$,

Convolution filter: $F[i, j]$

Convolution:

$$I[i, j] * F[i, j] = \sum_{n_1=-\infty}^{\infty} \sum_{n_2=-\infty}^{\infty} I[n_1][n_2] \cdot F[i - n_1][j - n_2]$$

Properties: translation equivariant

Example kernels 1

Operation	Kernel	Image result
Identity	$\begin{bmatrix} 0 & 0 & 0 \\ 0 & 1 & 0 \\ 0 & 0 & 0 \end{bmatrix}$	
Edge detection	$\begin{bmatrix} 1 & 0 & -1 \\ 0 & 0 & 0 \\ -1 & 0 & 1 \end{bmatrix}$	
Edge detection	$\begin{bmatrix} 0 & 1 & 0 \\ 1 & -4 & 1 \\ 0 & 1 & 0 \end{bmatrix}$	
Edge detection	$\begin{bmatrix} -1 & -1 & -1 \\ -1 & 8 & -1 \\ -1 & -1 & -1 \end{bmatrix}$	

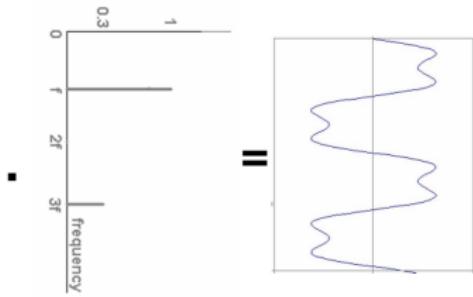
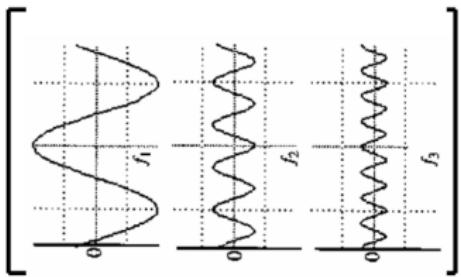
Example kernels 2

Sharpen	$\begin{bmatrix} 0 & -1 & 0 \\ -1 & 5 & -1 \\ 0 & -1 & 0 \end{bmatrix}$	
Box blur (normalized)	$\frac{1}{9} \begin{bmatrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{bmatrix}$	
Gaussian blur 3 × 3 (approximation)	$\frac{1}{16} \begin{bmatrix} 1 & 2 & 1 \\ 2 & 4 & 2 \\ 1 & 2 & 1 \end{bmatrix}$	
Gaussian blur 5 × 5 (approximation)	$\frac{1}{256} \begin{bmatrix} 1 & 4 & 6 & 4 & 1 \\ 4 & 16 & 24 & 16 & 4 \\ 6 & 24 & 36 & 24 & 6 \\ 4 & 16 & 24 & 16 & 4 \\ 1 & 4 & 6 & 4 & 1 \end{bmatrix}$	

Fourier transform

$$F(u) = \int_{-\infty}^{\infty} f(x) e^{-j2\pi ux} dx,$$

$$f(x) = \int_{-\infty}^{\infty} F(u) e^{j2\pi ux} du$$

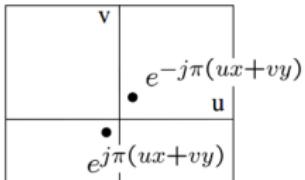


2D Fourier

$$F(u, v) = \int_{-\infty}^{\infty} \int_{-\infty}^{\infty} f(x, y) e^{-j2\pi(ux+vy)} dx dy,$$

$$f(x, y) = \int_{-\infty}^{\infty} \int_{-\infty}^{\infty} F(u, v) e^{j2\pi(ux+vy)} du dv$$

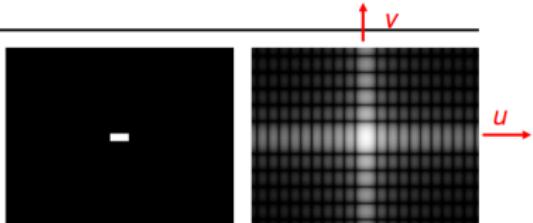
To get some sense of what basis elements look like, we plot a basis element --- or rather, its real part --- as a function of x, y for some fixed u, v . We get a function that is constant when $(ux+vy)$ is constant. The magnitude of the vector (u, v) gives a frequency, and its direction gives an orientation. The function is a sinusoid with this frequency along the direction, and constant perpendicular to the direction.



Fourier example 1

FT pair example 1

rectangle centred at origin
with sides of length X and Y



$$F(u, v) = \int \int f(x, y) e^{-j2\pi(ux+vy)} dx dy,$$

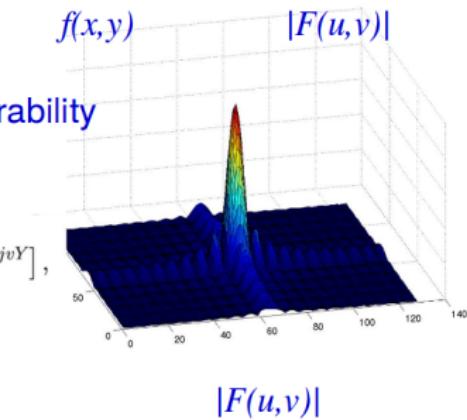
$$= \int_{-X/2}^{X/2} e^{-j2\pi ux} dx \int_{-Y/2}^{Y/2} e^{-j2\pi vy} dy, \text{ separability}$$

$$= \left[\frac{e^{-j2\pi ux}}{-j2\pi u} \right]_{-X/2}^{X/2} \left[\frac{e^{-j2\pi vy}}{-j2\pi v} \right]_{-Y/2}^{Y/2},$$

$$= \frac{1}{-j2\pi u} [e^{-juX} - e^{juX}] \frac{1}{-j2\pi v} [e^{-jvY} - e^{jvY}],$$

$$= XY \left[\frac{\sin(\pi Xu)}{\pi Xu} \right] \left[\frac{\sin(\pi Yv)}{\pi Yv} \right]$$

$$= XY \operatorname{sinc}(\pi Xu) \operatorname{sinc}(\pi Yv).$$



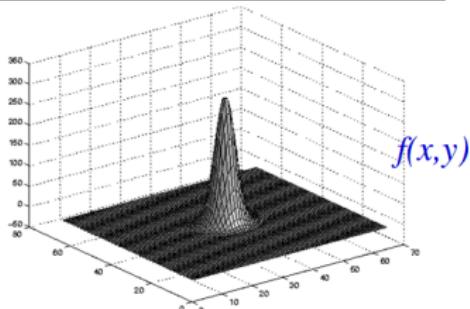
Fourier example 2

FT pair example 2

Gaussian centred on origin

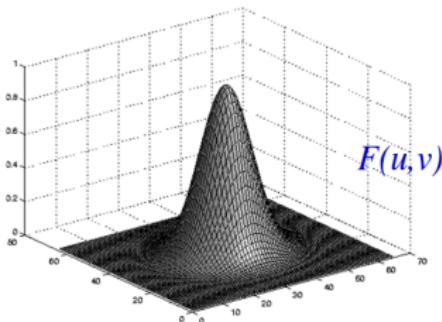
$$f(r) = \frac{1}{2\pi\sigma^2} e^{-r^2/2\sigma^2}$$

where $r^2 = x^2 + y^2$.



$$F(u, v) = F(\rho) = e^{-2\pi^2\rho^2\sigma^2}$$

where $\rho^2 = u^2 + v^2$.



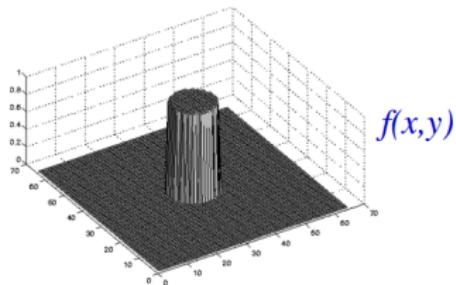
- FT of a Gaussian is a Gaussian
- Note inverse scale relation

Fourier example 3

FT pair example 3

Circular disk unit height and
radius a centred on origin

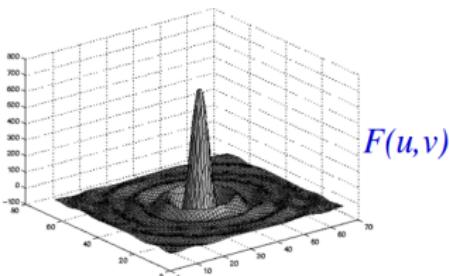
$$f(x, y) = \begin{cases} 1, & |r| < a, \\ 0, & |r| \geq a. \end{cases}$$



$$F(u, v) = F(\rho) = a J_1(\pi a \rho) / \rho$$

where $J_1(x)$ is a Bessel function.

- rotational symmetry
- a '2D' version of a sinc



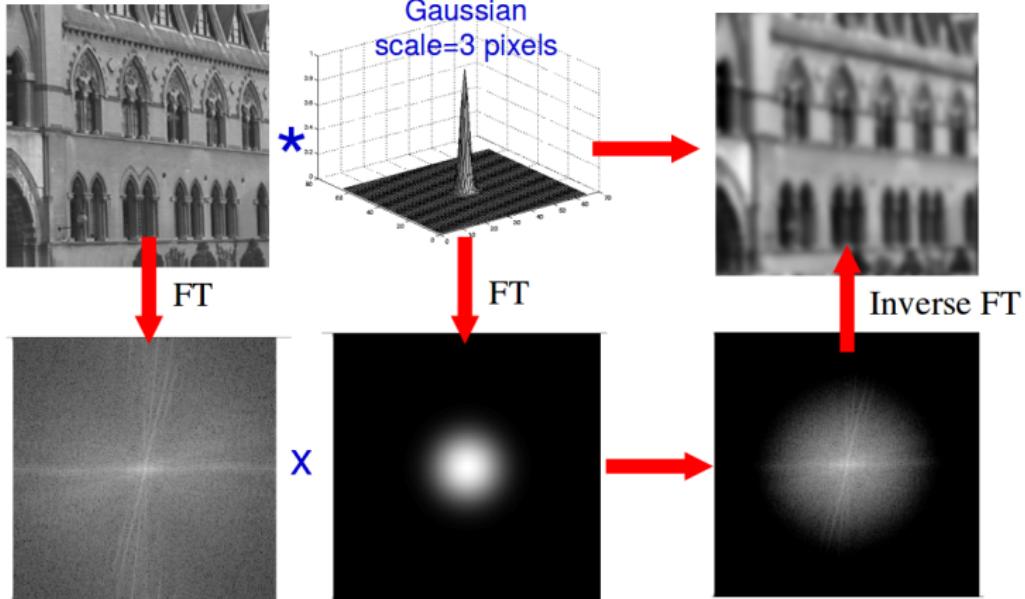
Phase importance



Fourier properties

- Fourier is a linear transform
- Fourier of convolution is equal to product of Fourier transforms
- Fourier of Gaussian is Gaussian
- FFT (Fast Fourier transform) computes it in $n \log(n)$

Example: image restoration



Blurring acts as a low pass filter and attenuates higher spatial frequencies

Problems with direct deblur

Deblurring with an inverse filter

noise $\sigma = 0.3$ grey levels

$$\hat{F}(u,v) = G(u,v) / H(u,v)$$

blur $\sigma = 0.5$ pixels



$g(x,y)$

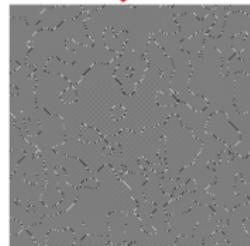
blur $\sigma = 1.0$ pixels



blur $\sigma = 1.5$ pixels



$\hat{f}(x,y)$



The Wiener filter

$$\hat{F}(u, v) = W(u, v) G(u, v)$$

$$W(u, v) = \frac{H^*(u, v)}{|H(u, v)|^2 + K(u, v)}$$

where

$$K(u, v) = S_\eta(u, v)/S_f(u, v)$$

$S_f(u, v) = |F(u, v)|^2$ power spectral density of $f(x, y)$

$S_\eta(u, v) = |N(u, v)|^2$ power spectral density of $\eta(x, y)$

Wiener deblur

Example 1: Focus deblurring with a Wiener filter

blur $\sigma = 1.5$ pixels

noise $\sigma = 0.3$ grey levels

$$\hat{F}(u, v) = W(u, v) G(u, v)$$

$$W(u, v) = \frac{H^*(u, v)}{|H(u, v)|^2 + K(u, v)}$$

$g(x, y)$



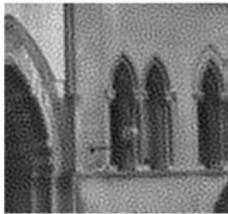
$\hat{f}(x, y)$



$K = 1.0 \text{ e } -5$

$K = 1.0 \text{ e } -3$

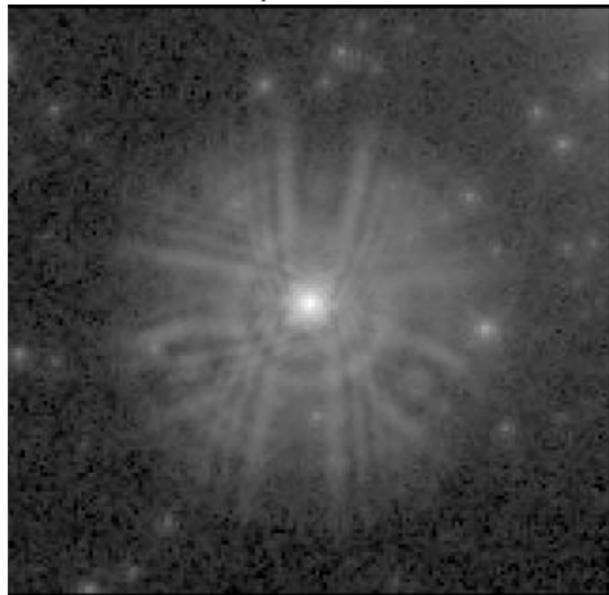
$K = 1.0 \text{ e } -1$



Point spread function

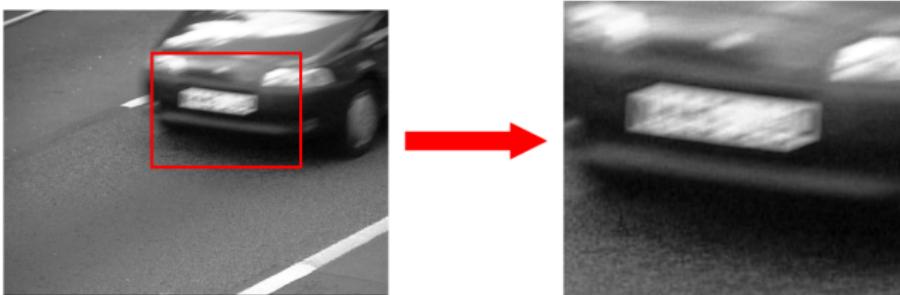
The point spread function (PSF) describes the response of an imaging system to a point source or point object.

Hubble telescope PSF:



Example: reading number plates

Application: Reading number plates



Algorithm

1. Rotate image so that blur is horizontal
2. Estimate length of blur
3. Construct a bar modelling the convolution
4. Compute and apply a Wiener filter
5. Optimize over values of K

Example: reading number plates

$f(x,y)$



$h(x,y)$



$\hat{f}(x,y)$



blur = 30 pixels

Optimization for deblurring

- Estimate $f(x,y)$ by optimizing a cost function:

$$\hat{f} = \arg \min_f \underbrace{(g - Af)^2}_{\text{Likelihood/loss function}} + \lambda \underbrace{p(f)}_{\text{prior/regularization}}$$

observed image generated image

Example

$$p(f) = (\nabla f)^2$$

to suppress high frequency noise

Blind deblurring

So far we have assumed that we know the generative model, e.g.

$$\mathbf{g} = \mathbf{A}(\mathbf{h}) \mathbf{f}$$

$$\mathbf{G} = \mathbf{H} \mathbf{F}$$



i.e. that $h(x,y)$ is known, so that given the observed image $g(x,y)$, then the original image $f(x,y)$ can be estimated (restored)

Consider if only the observed image $g(x,y)$ is known.
This is the problem of **blind estimation**.

Blind deblurring

- Estimate $f(x,y)$ and $h(x,y)$ by optimizing a cost function:

$$\min_{f,h} \underbrace{(g - A(h)f)^2}_{\text{Likelihood/loss function}} + \lambda p_f(f) + \mu p_h(h)$$

observed image generated image

↓

blur prior

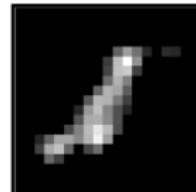
image prior

Blind deblurring

blurred image



estimated
blur filter



restored image



Blind deblurring



Sobel edge detector

Gradient in x and y direction:

$$\mathbf{G}_x = \begin{bmatrix} +1 & 0 & -1 \\ +2 & 0 & -2 \\ +1 & 0 & -1 \end{bmatrix} * \mathbf{A} \quad \text{and} \quad \mathbf{G}_y = \begin{bmatrix} +1 & +2 & +1 \\ 0 & 0 & 0 \\ -1 & -2 & -1 \end{bmatrix} * \mathbf{A}$$

Decomposition:

$$\begin{bmatrix} +1 & 0 & -1 \\ +2 & 0 & -2 \\ +1 & 0 & -1 \end{bmatrix} = \begin{bmatrix} 1 \\ 2 \\ 1 \end{bmatrix} [+1 \ 0 \ -1]$$

$$\mathbf{G} = \sqrt{\mathbf{G}_x^2 + \mathbf{G}_y^2}$$

$$\Theta = \text{atan}\left(\frac{\mathbf{G}_y}{\mathbf{G}_x}\right)$$

Sobel example



Sobel gradient



LoG: Laplacian of Gaussian

- To reduce the noise effect, the image is first smoothed.
- When the filter chosen is a Gaussian, we call it the LoG edge detector.

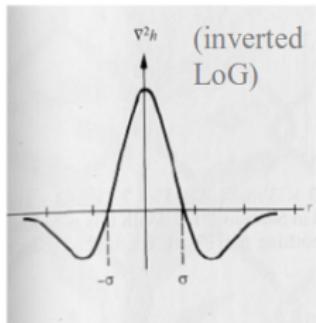
$$G(x, y) = e^{-\frac{x^2+y^2}{2\sigma^2}}$$

σ controls smoothing

- It can be shown that:

$$\nabla^2[f(x, y) * G(x, y)] = \nabla^2G(x, y) * f(x, y)$$

$$\nabla^2G(x, y) = \left(\frac{r^2 - 2\sigma^2}{\sigma^4}\right)e^{-r^2/2\sigma^2}, (r^2 = x^2 + y^2)$$



LoG: Laplacian of Gaussian

- The 2-D Laplacian of Gaussian (LoG) function centered on zero and with Gaussian standard deviation σ has the form:

$$LoG(x, y) = -\frac{1}{\pi\sigma^4} \left[1 - \frac{x^2 + y^2}{2\sigma^2}\right] e^{-\frac{x^2+y^2}{2\sigma^2}}$$

where σ is the standard deviation

- The amount of smoothing can be controlled by varying the value of the standard deviation.

LoG: Laplacian of Gaussian

(inverted LoG)

5 × 5 Laplacian of Gaussian mask

$$\begin{array}{ccccc} 0 & 0 & -1 & 0 & 0 \\ 0 & -1 & -2 & -1 & 0 \\ -1 & -2 & 16 & -2 & -1 \\ 0 & -1 & -2 & -1 & 0 \\ 0 & 0 & -1 & 0 & 0 \end{array}$$

(inverted LoG)

17 × 17 Laplacian of Gaussian mask

0	0	0	0	0	0	-1	-1	-1	-1	-1	0	0	0	0	0	0
0	0	0	-1	-1	-1	-1	-1	-1	-1	-1	0	0	0	0	0	0
0	0	-1	-1	-1	-2	-3	-3	-3	-3	-3	-2	-1	-1	-1	0	0
0	0	-1	-1	-2	-3	-3	-3	-3	-3	-3	-2	-2	-1	-1	0	0
0	-1	-1	-2	-3	-3	-3	-2	-3	-2	-3	-3	-3	-2	-1	-1	0
0	-1	-2	-3	-3	-3	0	2	4	2	0	-3	-3	-3	-2	-1	0
-1	-1	-3	-3	-3	0	4	10	12	10	4	0	-3	-3	-3	-1	-1
-1	-1	-3	-3	-2	2	10	18	21	18	10	2	-2	-3	-3	-1	-1
-1	-1	-3	-3	-3	4	12	21	24	21	12	4	-3	-3	-3	-1	-1
-1	-1	-3	-3	-2	2	10	18	21	18	10	2	-2	-3	-3	-1	-1
-1	-1	-3	-3	-3	0	4	10	12	10	4	0	-3	-3	-3	-1	-1
0	-1	-2	-3	-3	-3	0	2	4	2	0	-3	-3	-3	-2	-1	0
0	-1	-1	-2	-3	-3	-3	-2	-3	-2	-3	-3	-3	-2	-1	-1	0
0	0	-1	-1	-2	-3	-3	-3	-3	-3	-3	-2	-1	-1	-1	0	0
0	0	-1	-1	-1	-2	-3	-3	-3	-3	-3	-2	-1	-1	-1	0	0
0	0	0	0	-1	-1	-1	-1	-1	-1	-1	0	0	0	0	0	0

filtering



zero-crossings

Edge detection

Try out more edge detectors:

- Canny
- Deriche
- Prewitt
- Roberts cross